

# Evaluating Durability of Distributed Databases

## Theory and Empirical Studies of MongoDB

---

Kosta Dunn

Supervisor: Alan Fekete

Updated: 2018/11/12



Background



*Once a transaction has been committed, it will remain committed even in the case of a crash of 1 or more machines.*

It is one of the ACID properties, which ensures that an operation which has been acknowledged *cannot* be lost. (Haerder and Reuter, 1983)


Durability is violated if an *acknowledged write is lost*.

MongoDB is a widely used distributed database system. It is becoming one of the primary choices for storing critical user data.

- ⦿ Document model
- ⦿ Document = Row (in traditional database)
- ⦿ No schema


It provides high performance, availability and consistency. These properties are achieved via *replication*.

# Replication




Create a *replica set* by copying data across multiple machines...

# Replication



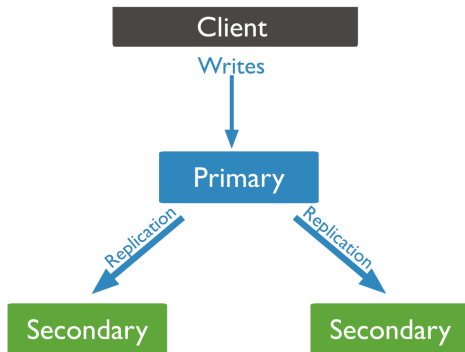
...So when one machine goes down...

# Replication



... The data (and hence the service) are still available!

# Replication in MongoDB

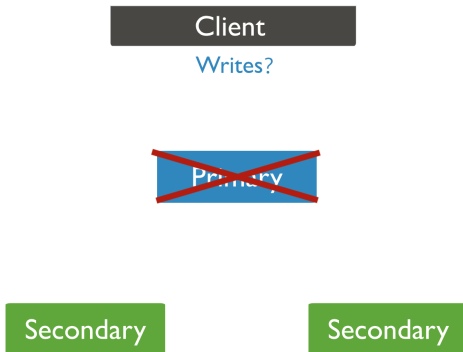


MongoDB uses a *Primary-Backup* strategy. **One** Primary, **the rest** are Secondaries.

All **write** operations must go through the Primary.




# Replication in MongoDB



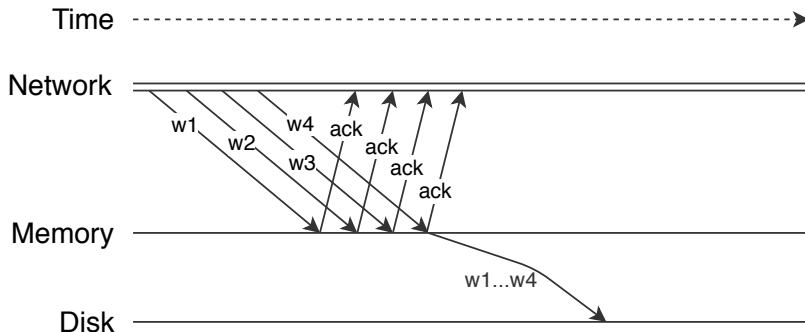
If a primary is down, **no write operations** can be acknowledged.

# Replication in MongoDB



When the Primary fails, all working Secondaries perform an *election* to select a new primary.

# MongoDB Journal



A journal is used to *recover* after failure by "replaying" the operations.

MongoDB with **Journal**ed writes will **send acknowledgements** without them going to disk!

MongoDB allows users to configure the *strength* of their writes through the **write concern**.

Stronger = harder to lose, but longer to acknowledge.

MongoDB allows users to configure the *strength* of their writes through the **write concern**.

Stronger = harder to lose, but longer to acknowledge.

**Primary** Operation is applied to the data on the primary

**Journal** Operation is applied to data and added to Journal on the primary

**Majority** Operation is applied and added to Journal on the *majority* of the replicas.

MongoDB allows users to configure the *strength* of their writes through the **write concern**.

Stronger = harder to lose, but longer to acknowledge.

**Primary** Operation is applied to the data on the primary

**Journal** Operation is applied to data and added to Journal on the primary

**Majority** Operation is applied and added to Journal on the *majority* of the replicas.

Under **Majority** write concern, there should be *no write loss!*

# What has been done

Brewer (2000) conjectured that databases cannot be simultaneously *Consistent, Available and Network failure tolerant*. Gilbert and Lynch (2002) proved this result.

⇒ There are limits to the capability of any database!

A tool "Jepsen" was developed to test effects of *Network failures* on databases. (Kingsbury, 2013, 2017; Patella, 2018)

And Alagappan et al. (2016) studied durability when *all* replicas fail, with a focus on file systems.

Note: probability of *all* replicas failing is **incredibly slim**.

There is **no work** on durability of distributed databases under single machine failures.

*To equip users and designers of distributed databases with the means to protect their systems from durability failures.*



- ⊙ **Categorisation of scenarios that result in write loss.**
- ⊙ **Design of an experiment capable of inducing write loss.**
- ⊙ Algorithm to quantify the number of lost writes.
- ⊙ **Empirical results to show that the experiment and algorithm work by detecting bugs in MongoDB 3.6-rc0.**
- ⊙ Theoretical model for evaluating when a write becomes durable.
- ⊙ **Estimation of when a write becomes durable on the Primary, using rudimentary client-accessible measurements.**
- ⊙ **An empirical study of time-till-durability on the Primary for acknowledged writes.**

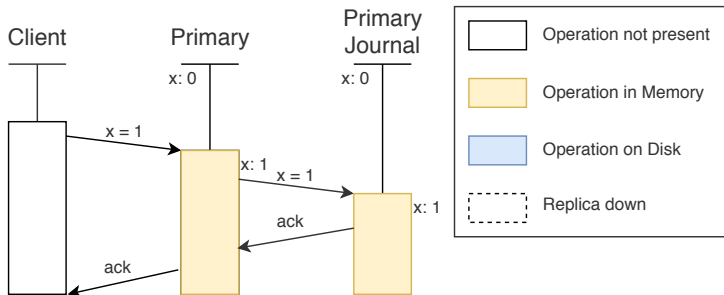
# Detecting and Quantifying Durability Failures

---

Can we create a scenario that *forces* MongoDB to lose a write?

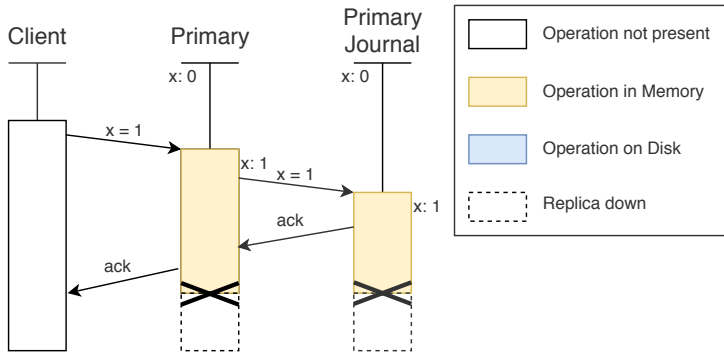
How big is the impact of these scenarios?

# Write Loss scenario



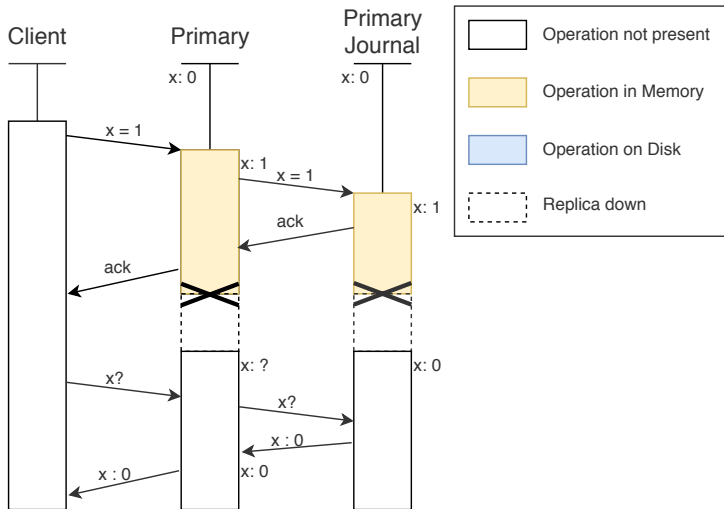
We issue a **Journaled** write and get an *acknowledgement*.

# Write Loss scenario



The Primary *crashes*.

# Write Loss scenario



The Primary *loses* any writes still *in-memory*.

# How do we do this empirically?

We created a tool that:

- ⦿ Configures a replica set
- ⦿ Stresses it with reads and writes
- ⦿ Crashes the Primary **100 seconds** into experiment.
- ⦿ Recovers the (old) Primary **200 seconds** into experiment.
- ⦿ Observes *incorrect values* in queries.

# Results - Write Loss: Mongo 4.0

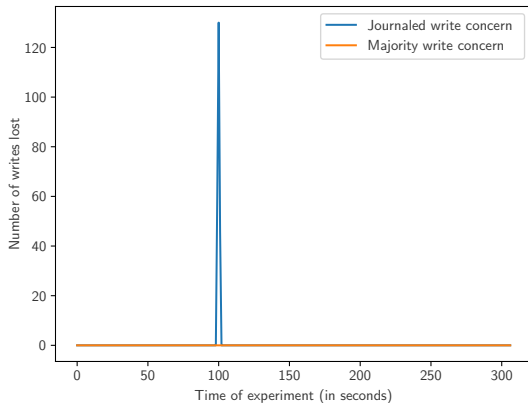


Figure: Distribution of write loss for every second in MongoDB 4.0.



# Results - Write Loss: Mongo 3.6-rc0




Figure: Distribution of write loss for every second in MongoDB 3.6-rc0.

MongoDB 4.0 performs as predicted by our theory. Found losses where we expect based on our theory.

MongoDB 3.6-rc0 loses writes where they *shouldn't be lost*.

⇒ Our tool succeeded in detecting bugs in MongoDB 3.6-rc0.

Our tool works!

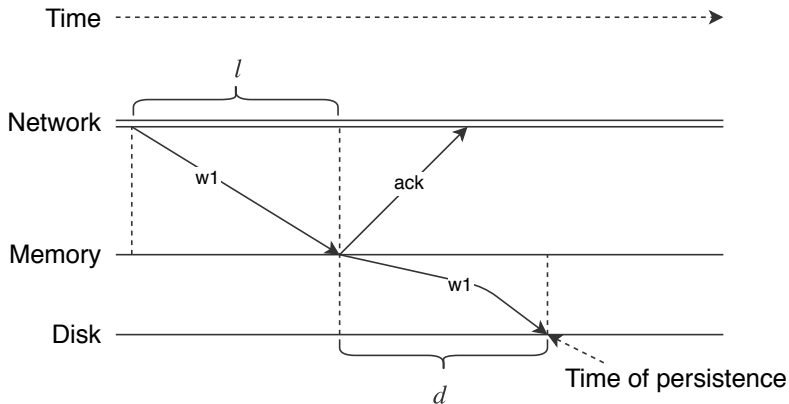
## Estimating when Writes Become Durable

---

Write loss is more common the closer the write is to the failure.

Can we estimate when our writes becomes durable, at least on the *Primary* replica?


# When does a write become durable?



A write becomes *durable* when it gets *persisted to disk*.

In other words  $l + d$ .

# Problem



**We can't measure  $d$ .**

# What can we measure?

- ⊙ Latency / (ping)
- ⊙ Response time of write operations...



# What can we measure?

- ⊙ Latency / (ping)
- ⊙ Response time of write operations... With different **write concerns**

# Estimating $d$ from Write Concerns

Which **write concern** persists a write to *disk*?

# Estimating $d$ from Write Concerns


Which **write concern** persists a write to *disk*?

**Journalled!**

A **Journalled** write will be *applied to memory* and *added to the journal* before being acknowledged.

Since the journal does not always hit the disk, we have only an *approximation*.

# Modelling Journaled Write Concern



Behaviour of a **Journaled** write.

That means, a **Journaled** write is:

$$j = l + d_{est} + l$$

# Estimating durability

We want to estimate  $t = l + d$  (time to durability).

We know  $l$  and  $j = l + d_{est} + l$ .

We then define the **estimate** as  $t_{est} = j - l$ :

$$\begin{aligned} t_{est} &= j - l \\ &= l + d_{est} \\ &\approx l + d \end{aligned}$$

# Results - latency distribution

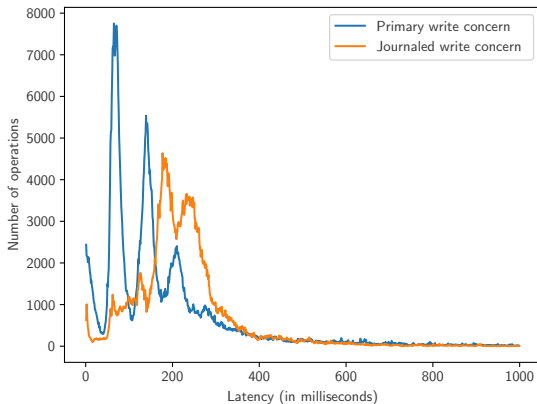


Figure: A frequency graph of the number of write operations acknowledged at latencies of 1-1000ms.



# Results - cumulative distribution

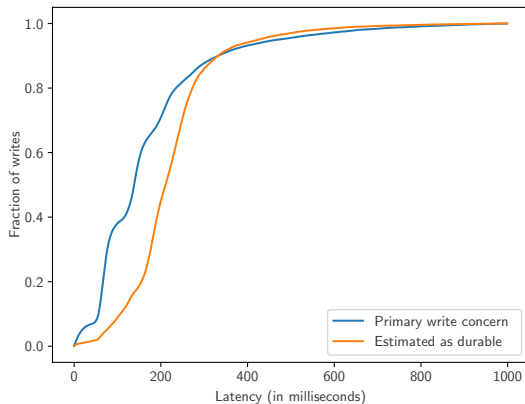


Figure: A cumulative frequency graph of the writes which become acknowledged and 1-durable at latencies 1-1000ms.

We can *estimate* when a write becomes durable on the primary by looking *how long* a **Journaled** write takes to come back.

In our case, 90% of writes are durable by 300ms.

## Conclusion

---

# Thesis Findings

- ⦿ Identified and categorised scenarios which cause writes to be lost.
- ⦿ Created a tool capable of inducing write loss.
- ⦿ Designed an algorithm to quantify write loss.
- ⦿ Showed that our tool works by finding bugs in MongoDB.
- ⦿ Derived a formula for when a write becomes durable on any number of replicas.
- ⦿ Developed an estimation of when a write becomes durable on the *Primary*, using rudimentary client-accessible measurements.
- ⦿ Found that 90% of writes are durable by 300ms after submission.

- ⦿ Induced only **one** failure per experiment, only on the Primary replica.
- ⦿ Only used client-accessible measurements.
- ⦿ Focused on durability only on the Primary.

- ⦿ Induce multiple failures, on any replica
- ⦿ Investigate MongoDB's own logs
- ⦿ Explore estimating durability on any number of replicas

This presentation focused on a subset of our thesis contributions. Here is the complete list:

- ⦿ Categorisation of *scenarios* that result in write loss.
- ⦿ Design of an experiment capable of inducing write loss.
- ⦿ Algorithm to quantify the number of lost writes.
- ⦿ Empirical results to show that the experiment and algorithm work by detecting bugs in MongoDB 3.6-rc0.
- ⦿ Theoretical model for evaluating when a write becomes durable.
- ⦿ Estimation of when a write becomes durable on the *Primary*, using rudimentary client-accessible measurements.
- ⦿ An empirical study of time-till-durability on the *Primary* for acknowledged writes.

Ramnatthan Alagappan, Aishwarya Ganesan, Yuvraj Patel, Thanumalayan Sankaranarayanan Pillai, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. Correlated crash vulnerabilities. In *Proceedings of the 12th USENIX Conference on Operating Systems Design and Implementation*, OSDI'16, pages 151–167, Berkeley, CA, USA, 2016. USENIX Association. ISBN 978-1-931971-33-1. URL <http://dl.acm.org/citation.cfm?id=3026877.3026890>.



- Eric A Brewer. Towards Robust Distributed Systems. In *Proceedings of the Nineteenth Annual ACM Symposium on Principles of Distributed Computing*, PODC '00, pages 7—, New York, NY, USA, 2000. ACM. ISBN 1-58113-183-6. doi: 10.1145/343477.343502. URL <http://doi.acm.org.ezproxy1.library.usyd.edu.au/10.1145/343477.343502>.
- Seth Gilbert and Nancy Lynch. Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services. *ACM SIGACT News*, 33(2):51, jun 2002. ISSN 01635700. doi: 10.1145/564585.564601. URL <http://portal.acm.org/citation.cfm?doid=564585.564601>.

Theo Haerder and Andreas Reuter. Principles of transaction-oriented database recovery. *ACM Computing Surveys*, 15(4):287–317, dec 1983. ISSN 03600300. doi: 10.1145/289.291. URL <http://portal.acm.org/citation.cfm?doid=289.291>.

Kyle Kingsbury. Jepsen: MongoDB <https://aphyr.com/posts/284-call-me-maybe-mongodb>, 2013. URL <https://aphyr.com/posts/284-call-me-maybe-mongodb>.

Kyle Kingsbury. MongoDB 3.4.0-rc3 <https://jepsen.io/analyses/mongodb-3-4-0-rc3>, 2017. URL <https://jepsen.io/analyses/mongodb-3-4-0-rc3>.

Kit Patella. MongoDB 3.6.4

<https://jepsen.io/analyses/mongodb-3-6-4>, 2018. URL

<https://jepsen.io/analyses/mongodb-3-6-4>.